

CSC384h: Intro to Artificial Intelligence

• Knowledge Representation

- This material is covered in chapters 7—9 and 12 of the text.
- Chapter 7 provides a useful motivation for logic, and an introduction to some basic ideas. It also introduces propositional logic, which is a good background for first-order logic.
- What we cover here is mainly covered in Chapters 8 and 9. However, Chapter 8 contains some additional useful examples of how first-order knowledge bases can be constructed. Chapter 9 covers forward and backward chaining mechanisms for inference, while here we concentrate on resolution.
- Chapter 12 covers some of the additional notions that have to be dealt with when using knowledge representation in AI.

Knowledge Representation

- What is **knowledge**?
- **Information** we have about the world we inhabit
 - Both the physical and mental world.
 - We have knowledge about many abstract mental constructs and ideas
- Besides knowledge we have various other mental attitudes and feelings about our environment.
 - John knows “...”
 - John fears “...”
 - Then things get complex: John knows that he fears “...”
 - So knowledge can take a variety of forms, some quite complex

Knowledge Representation

- What is **Representation**?
- Symbols standing for things in the world



CSC 384

Words we use
in language

Symbols we use
in mathematics

- Can all knowledge be symbolically represented?

Knowledge Representation

- Can all knowledge be symbolically represented?
- No - we do not symbolically represent the “pixels” that we perceive at the back of our retina.
- So intelligent agents also perform a great deal of low level “non-symbolic reasoning” over their perceptual inputs.
- But higher level “symbolically represented” knowledge also seems to be essential
 - This is the kind of knowledge that we learn in school, by reading, etc.
- In this module we study symbolically represented knowledge

Reasoning

- What is reasoning (in the context of symbolically represented knowledge) ?
 - Manipulating our symbols to produce new symbols that represent new knowledge.

Deriving a new sentence

- Typically “symbols” are sequences of symbols, e.g., words in language sequenced together to form sentences.
- So we will develop methods for manipulating “sentences” to produce new “sentences”

Reasoning

- In language we can make up a huge variety of sentences.
- Each of these sentences makes some sort of claim or assertion about our world (mental or physical).
- These claims could be true or false.
 - I am anxious, so the sentence “I feel calm and relaxed.” Is false
- Reasoning aims to be **TRUTH PRESERVING**.
- If we use reasoning to manipulate a collection of **TRUE** sentences, we want the newly derived sentences to also be **TRUE**
- If our reasoning is truth preserving we say that is is **SOUND**

Reasoning

- A more subtle idea is **COMPLETENESS**.
- Completeness says that our reasoning system is powerful enough to produce **ALL** sentences that must be true given one current collection of true sentences.
- Completeness requires a formal characterization of “sentence” in order to answer the question of if we have produced **ALL** true sentences

Knowledge Representation

- Consider the task of understanding a simple story.
- How do we test understanding?
- Not easy, but understanding at least entails some ability to answer simple questions about the story.

Example.

- ▶ Three little pigs



Example.

- ▶ Three little pigs



Example.

- Why couldn't the wolf blow down the house made of bricks?
- What background knowledge are we applying to come to that conclusion?

Why Knowledge Representation?

- Large amounts of knowledge are used to understand the world around us, and to communicate with others.
- We also have to be able to reason with that knowledge.
 - Our knowledge won't be about the blowing ability of wolfs in particular, it is about physical limits of objects in general.
 - We have to employ reasoning to make conclusions about the wolf.
 - More generally, reasoning **provides an exponential or more compression in the knowledge we need to store**. I.e., without reasoning we would have to store a infeasible amount of information: e.g., Elephants can't fit into teacups.

Logical Representations

- AI typically employs logical representations of knowledge.
- Logical representations useful for a number of reasons:

Logical Representations

- They are mathematically precise, thus we can analyze their limitations, their properties, the complexity of inference etc.
- They are formal languages, thus computer programs can manipulate sentences in the language.
- They come with both a formal **syntax** and a formal **semantics**.
- Typically, have well developed **proof theories**: formal procedures for reasoning at the syntactic level (achieved by manipulating sentences).
- In this module we will study First-Order logic, and a reasoning mechanism called resolution that operates on First-Order logic.

First Order Logic (FOL)

- Two components: Syntax and Semantics.
 - In a programming language we have a syntax for an if statement: “if <boolean condition>:<expressions>”
 - The if statement also has semantics: if <boolean condition> evaluates to TRUE then we execute <expressions>.
- Syntax gives the grammar or rules for forming proper sentences.
- Semantics gives the meaning.

Basic Semantic entities of FOL

- We have a set of objects **D**. These are objects in the world that are important for our application.
 - Often we will want to form tuples of objects, e.g., (d_1, d_2) where $d_1 \in \mathbf{D}$ and $d_2 \in \mathbf{D}$ are a pair of objects
 - A k-ary tuple is a subset of $\mathbf{D}^k = \mathbf{D} \times \mathbf{D} \times \dots \times \mathbf{D}$ the k-wise Cartesian product of **D**
- We can identify special sets of objects (subsets of **D**) that have some property in common. These sets are called **properties (predicates)**.
 - E.g., **female**, **male**, **children**, **adult** could each need subsets that we identify as being useful in our application. If an object **d** is in the set **male**, we can say that d has the property male : **male(d)**.

Basic Semantic entities of FOL

- Sometimes individual objects are not sufficient, we want to identify special groups (tuples) of objects that are related to each other. We call these sets **relations**.
 - E.g. **married** might be a special subset of pairs that we wish to keep track of in our application.
- Finally, we might want to keep track of functions over our objects. $f: \mathbf{D} \rightarrow \mathbf{D}$
 - E.g. for $d \in \text{student}$, we might want a function $\text{faculty}(d)$ that gives the faculty the student is registered in.
 - More generally we might want $f: \mathbf{D}^k \rightarrow \mathbf{D}$, i.e., a function of many arguments mapping \mathbf{D} .

Basic Syntactic symbols of FOL

- The syntax starts off with a different symbol for each basic semantic entity (objects, functions, predicates, relations) that we have decided to utilize.
- We get to decide what symbols we use (but of course want to use symbols that are easy to understand)
- These user specified symbols are called the **primitive symbols**

Syntax	Semantics
Constant symbols	A particular object $d \in \mathbf{D}$
Function symbols	Some function $f: \mathbf{D}^k \rightarrow \mathbf{D}$
Predicate symbols	Some subset of \mathbf{D}
Relation symbols	Some subset of \mathbf{D}^k

Basic Syntactic symbols of FOL

- In addition we introduce some additional symbols that we will use to connect our basic symbols into sentences.

Syntax	Semantics
Constant symbols	A particular object $d \in \mathbf{D}$
Function symbols	Some function $f: \mathbf{D}^k \rightarrow \mathbf{D}$
Predicate symbols	Some subset of \mathbf{D}
Relation symbols	Some subset of \mathbf{D}^k
Equality (commonly used relation)	Subset of $\mathbf{D}^2 = \{(d,d) \mid d \in \mathbf{D}\}$
Variables (as many as we need) x, y, z, \dots	An object $d \in \mathbf{D}$ (which particular object can vary)
Logical connectives: $\wedge, \vee, \neg, \rightarrow$...defined below...
Quantifiers: \forall, \exists	...defined below...

Example

- Teaching CSC384, want to represent knowledge that would be useful for making the course a successful learning experience. So we might choose syntactic symbols like
- Objects:
 - students, subjects, assignments, numbers.
- Predicates:
 - $\text{difficult}(\textit{subject})$, $\text{CSMajor}(\textit{student})$
- Relations:
 - $\text{handedIn}(\textit{student}, \textit{assignment})$
- Functions:
 - $\text{Grade}(\textit{student}, \textit{assignment}) \rightarrow \textit{number}$

First Order Syntax (the grammar)

- We start with out basic syntactic symbols **constants**, **functions**, **predicates**, **relations** and **variables**.
 - Note: the function and relation symbols each have specific arities (the number of arguments it takes)
- From these we can build upon **terms** and **sentences(formulas)**. Terms are ways of applying functions to build up new “names” for objects. Formulas, are denoting true/false assertions about terms.

First Order Syntax - Terms

- Terms are used as names (perhaps complex nested names) for objects in the domain.

Terms	
Constants	$c, \text{john}, \text{mary}$
Variables	x, y, z, \dots
Function application	$f(t_1, t_2, \dots, t_k)$ t_i are already constructed terms

- 5 is a constant term: a symbol representing the number 5.
john is a term — a symbol representing the person John.
- $+(5,5)$ is a function application term — a new symbol representing the number 10.

First Order Syntax - Terms

- **Note:** constants are the same as functions taking zero arguments.
- Terms are names for objects (things in the world):
 - Constants denote specific objects
 - Functions map tuples of objects to other objects
 - bill, jane, father(jane), father(father(jane))
 - **X**, father(**X**), hotel7, rating(hotel7), cost(hotel7)
 - Variables like X are not yet determined, but they will eventually denote particular objects.

First Order Syntax - Sentences.

- Once we have terms we can build up *sentences (formulas)*
Terms represent objects, *formulas* represent true/false assertions about these objects

First Order Syntax - Sentences.

Formula	
Atomic formula	$p(t)$ or $r(t_1, t_2, \dots, t_k)$ p is a predicate symbol, r is a k-ary relation symbol, t_i are terms
Negation	$\neg f$ f is a fomula
Conjunction	$f_1 \wedge f_2 \wedge \dots \wedge f_k$ f_i are formulas
Disjunction	$f_1 \vee f_2 \vee \dots \vee f_k$
Implication	$f_1 \rightarrow f_2$ f_1 and f_2 are fomulas f_1 often calles the antecedent, f_2 the consequence
Existential	$\exists X.f$ f is a formula X is a variable
Universal	$\forall X.f$

Intuition (formalized later).

- Atoms denote facts that can be true or false about the world
 - father_of(jane,bill), female(jane), system_down()
 - satisfied(client15), **satisfied(C)**
 - desires(client15,rome,week29), **desires(X,Y,Z)**
 - rating(hotel7,4), cost(hotel7,125)
- Other formulas generate more complex assertions by composing these atomic formulas.
 - Their truth is dependent on the truth of the atomic formulas in them

Semantics

- Formulas (syntax) can be built up recursively, and can become arbitrarily complex
- Intuitively, there are various distinct formulas (viewed as strings) that really are asserting the same thing
 - $\forall X, Y. \text{elephant}(X) \wedge \text{teacup}(Y) \rightarrow \text{largerThan}(X, Y)$
 - $\forall X, Y. \text{teacup}(Y) \wedge \text{elephant}(X) \rightarrow \text{largerThan}(X, Y)$
- To capture this equivalence and to make sense of complex formulas we utilize the semantics

Semantics

- A formal mapping from formulas to true/false assertions about our semantic entities (individuals, sets and relations over individuals, functions over individuals).
- The mapping mirrors the recursive structure of the syntax, so we can map any formula to a composition of assertions about the semantic entities.

Syntax - The language

- First, we must fix the particular first-order language we are going to provide semantics for. The **primitive** symbols included in the syntax defines the particular language.

$L(F,P,V)$

F = set of function (and constant symbols)

Each symbol f in F has a particular arity.

P = set of predicate and relation symbols.

Each relation symbol $r \in P$ has a particular arity. (The predicate symbols always have arity 1)

V = an infinite set of variables.

Semantics - Primitive Symbols

- An **interpretation** (model) specifies the mapping from the primitive symbols to semantic entities. It is a tuple $\langle D, \Phi, \Psi, V \rangle$
 - D is a non-empty set of objects (domain of discourse)
 - Φ specifies the meaning of each primitive function symbol
 - Also handles the primitive constant symbols (these can be viewed as being zero-arity functions).
 - Ψ specifies the meaning of each primitive predicate and relation symbol.
 - V specifies the meaning of the variables.
- Note, the semantic entities that a syntactic symbol maps to is often called the **meaning** of the symbol or the **denotation** of the symbol

Semantics - Primitive Symbols

Symbol	Semantics
Constant Symbol c	$\Phi(c) \in \mathbf{D}$ (some particular object)
K-ary function symbol f	$\Phi(f)$ Some particular function $\mathbf{D}^k \rightarrow \mathbf{D}$
Predicate symbol p	$\Psi(p)$ Some particular subset of \mathbf{D}
K-ary relation symbol r	$\Psi(r)$ Some particular subset of \mathbf{D}^k
Variable x	$V(x) \in \mathbf{D}$ (some particular object)

Intuitions: Domain

- Domain D : $d \in D$ is an *individual*
- E.g. $\{\underline{\text{craig}}, \underline{\text{jane}}, \underline{\text{grandhotel}}, \underline{\text{marriot}}, \underline{\text{rome}}, \underline{\text{portofino}}, \underline{100}, \underline{110}, \underline{120} \dots\}$
- We use underlined symbols to talk about domain individuals (syntactic symbols of the first-order language are not underlined)
- Domains often infinite, but we'll use finite models to prime our intuitions

Intuitions: Φ

- Given k -ary function f and k individuals $d_1 \dots d_k$, what individual does $f(d_1, \dots, d_k)$ denote
 - Constants (0-ary functions) are mapped to individuals in \mathbf{D} .
 - $\Phi(\text{client17}) = \underline{\text{craig}}$, $\Phi(\text{hotel5}) = \underline{\text{marriot}}$, $\Phi(\text{rome}) = \underline{\text{rome}}$
 - 1-ary functions are mapped to particular functions in $\mathbf{D} \rightarrow \mathbf{D}$
 - $\Phi(\text{rating}) = \underline{\text{f_rating}}$:
 - $\underline{\text{f_rating}}(\text{grandhotel}) = \underline{\text{5stars}}$
 - 2-ary functions are mapped to functions from $\mathbf{D}^2 \rightarrow \mathbf{D}$
 - $\Phi(\text{distance}) = \underline{\text{f_distance}}$:
 - $\underline{\text{f_distance}}(\text{toronto}, \text{sienna}) = \underline{\text{3256}}$
 - N-ary functions are mapped similarly

Intuitions: Ψ

- Given k-ary relation r , what does r denote
- 0-ary predicates are mapped to true or false.
 $\Psi(\text{rainy}) = \text{True}$ $\Psi(\text{sunny}) = \text{False}$
- 1-ary predicates are mapped to subsets of \mathbf{D} .
 - $\Psi(\text{privatebeach}) = p_privatebeach$: (the subset of hotels that have a private beach)
e.g. $p_privatebeach = \{\underline{grandhotel}, \underline{fourseasons}\}$
- 2-ary predicates are mapped to subsets of \mathbf{D}^2 (sets of pairs of individuals)
 - $\Psi(\text{location}) = p_location$: $p_location(\underline{grandhotel}, \underline{rome}) = \text{True}$
 $p_location(\underline{grandhotel}, \underline{sienna}) = \text{False}$
 - $\Psi(\text{available}) = p_available$:
 $p_available(\underline{grandhotel}, \underline{week29}) = \text{True}$
- n-ary predicates..subsets of \mathbf{D}^n

Intuitions: v

- V exists to take care of quantification. As we will see the exact mapping it specifies will not matter.

Semantics — Terms

- Given language $L(F,P,V)$, and an interpretation $\mathbf{I} = \langle \mathbf{D}, \Phi, \Psi, V \rangle$ and a term \mathbf{t} . $\mathbf{I}(\mathbf{t})$ is the denotation of \mathbf{t} under \mathbf{I} .

Term	Semantics
Constant Symbol \mathbf{c}	$\mathbf{I}(\mathbf{c}) = \Phi(\mathbf{c}) \in \mathbf{D}$ (some particular object)
Variable \mathbf{x}	$\mathbf{I}(\mathbf{x}) = V(\mathbf{x}) \in \mathbf{D}$ (some particular object)
Function application $\mathbf{f}(\mathbf{t}_1, \mathbf{t}_2, \dots, \mathbf{t}_k)$	$\mathbf{I}(\mathbf{f}(\mathbf{t}_1, \mathbf{t}_2, \dots, \mathbf{t}_k)) = \Phi(\mathbf{f})(\mathbf{I}(\mathbf{t}_1), \mathbf{I}(\mathbf{t}_2), \dots, \mathbf{I}(\mathbf{t}_k))$ First we obtain the denotation of each argument under \mathbf{I} , then we apply the function $\Phi(\mathbf{f})$ to these interpreted terms

- Hence the terms always denote individuals under interpretation \mathbf{I}

Semantics — Formulas

- Formulas will always be True or False under any interpretation I .

Formula	Semantics
Atomic formula $r(t_1, t_2, \dots, t_k)$	$I(r(t_1, t_2, \dots, t_k)) =$ True if $(I(t_1), I(t_2), \dots, I(t_k)) \in \Psi(r)$ False otherwise First we obtain the denotation of each argument under I . Then we check if this tuple of interpreted terms is in the set of tuples $\Psi(r)$

- Ψ Maps r to a subset of D^k (a subset of k -ary tuples of individuals). So the atomic formula is true if its arguments are in the stated relation.

Semantics — Formulas

Formula	Semantics
$\neg f$	$I(\neg f) =$ True if $I(f) = \text{False}$ False otherwise
$f_1 \wedge f_2 \wedge \dots \wedge f_k$	$I(f_1 \wedge f_2 \wedge \dots \wedge f_k) =$ True if $I(f_i) = \text{True}$ for every i False otherwise
$f_1 \vee f_2 \vee \dots \vee f_k$	$I(f_1 \vee f_2 \vee \dots \vee f_k) =$ True if $I(f_i) = \text{True}$ for any i False otherwise
$f_1 \rightarrow f_2$	$I(f_1 \rightarrow f_2) =$ True if $I(f_1) = \text{False}$ or $I(f_2) = \text{True}$ False otherwise

- Standard rules for proposition logic that you would have seen before (check chap 7 if not)

Semantics — Formulas

Formula	Semantics
$\exists X.f$	$I(f) =$ True if for some $d \in \mathbf{D}$, $I'(f) = \text{True}$ $I' = \langle \mathbf{D}, \Phi, \Psi, V[X=d] \rangle$ False Otherwise
$\forall X.f$	$I(f) =$ True if for all $d \in \mathbf{D}$, $I'(f) = \text{True}$ $I' = \langle \mathbf{D}, \Phi, \Psi, V[X=d] \rangle$ False Otherwise

- Quantifiers. Exists checks if f is true under some different variable mapping for the variable X . Forall checks if f is true under all possible mappings of the variable X .

Example

$D = \{\underline{\text{bob}}, \underline{\text{jack}}, \underline{\text{fred}}\}$

$I(\text{happy} = \{\underline{\text{bob}}, \underline{\text{jack}}, \underline{\text{fred}}\})$

$I(\forall X. \text{happy}(X))$

1. $\Psi(\text{happy})(X), (\forall[X = \text{bob}]) = \Psi(\text{happy})(\text{bob}) = \text{True}$

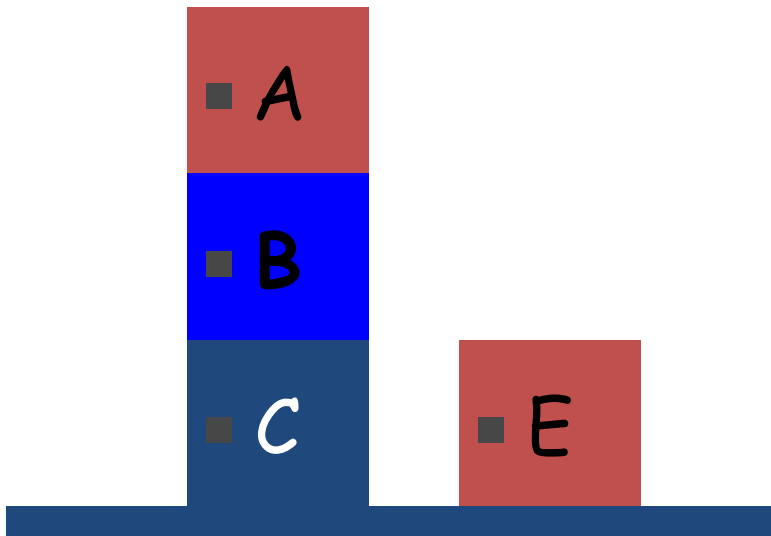
2. $\Psi(\text{happy})(X), (\forall[X = \text{jack}]) = \Psi(\text{happy})(\text{jack}) = \text{True}$

3. $\Psi(\text{happy})(X), (\forall[X = \text{fred}]) = \Psi(\text{happy})(\text{fred}) = \text{True}$

Therefore $I(\forall X. \text{happy}(X)) = \text{True}$.

Models—Examples.

Environment



Language (Syntax)

- Constants: a,b,c,e
- Functions:
 - No function
- Predicates:
 - on: binary
 - above: binary
 - clear: unary
 - ontable: unary

Models—Examples.

Language (syntax)

- Constants: a,b,c,e
- Predicates:
 - on (binary)
 - above (binary)
 - clear (unary)
 - ontable(unary)

A possible Model I_1 (semantics)

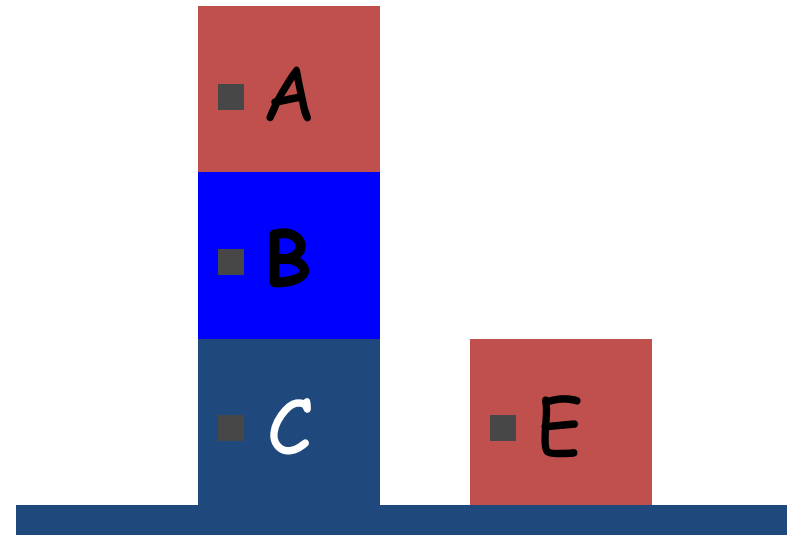
- $D = \{\underline{A}, \underline{B}, \underline{C}, \underline{E}\}$
- $\Phi(a) = \underline{A}, \Phi(b) = \underline{B}, \Phi(c) = \underline{C}, \Phi(e) = \underline{E}.$
- $\Psi(\text{on}) = \{(\underline{A}, \underline{B}), (\underline{B}, \underline{C})\}$
- $\Psi(\text{above}) = \{(\underline{A}, \underline{B}), (\underline{B}, \underline{C}), (\underline{A}, \underline{C})\}$
- $\Psi(\text{clear}) = \{\underline{A}, \underline{E}\}$
- $\Psi(\text{ontable}) = \{\underline{C}, \underline{E}\}$

Models—Examples.

Model I_1

- $D = \{\underline{A}, \underline{B}, \underline{C}, \underline{E}\}$
- $\Phi(a) = \underline{A}, \Phi(b) = \underline{B}, \Phi(c) = \underline{C}, \Phi(e) = \underline{E}.$
- $\Psi(\text{on}) = \{(\underline{A}, \underline{B}), (\underline{B}, \underline{C})\}$
- $\Psi(\text{above}) = \{(\underline{A}, \underline{B}), (\underline{B}, \underline{C}), (\underline{A}, \underline{C})\}$
- $\Psi(\text{clear}) = \{\underline{A}, \underline{E}\}$
- $\Psi(\text{ontable}) = \{\underline{C}, \underline{E}\}$

Environment



Models—Formulas true or false?

Model I_1

- $D = \{\underline{A}, \underline{B}, \underline{C}, \underline{E}\}$
- $\Phi(a) = \underline{A}, \Phi(b) = \underline{B}, \Phi(c) = \underline{C}, \Phi(e) = \underline{E}.$
- $\Psi(\text{on}) = \{(\underline{A}, \underline{B}), (\underline{B}, \underline{C})\}$
- $\Psi(\text{above}) = \{(\underline{A}, \underline{B}), (\underline{B}, \underline{C}), (\underline{A}, \underline{C})\}$
- $\Psi(\text{clear}) = \{\underline{A}, \underline{E}\}$
- $\Psi(\text{ontable}) = \{\underline{C}, \underline{E}\}$

$\forall X, Y. \text{on}(X, Y) \rightarrow \text{above}(X, Y)$

$X = \underline{A}, Y = \underline{B}. \quad ?$

$X = \underline{C}, Y = \underline{A} \quad ?$

...

$\forall X, Y. \text{above}(X, Y) \rightarrow \text{on}(X, Y)$

$X = \underline{A}, Y = \underline{B} \quad ?$

$X = \underline{A}, Y = \underline{A} \quad ?$

$X = \underline{A}, Y = \underline{C} \quad ?$

Models—Examples.

Model I_1

- $D = \{\underline{A}, \underline{B}, \underline{C}, \underline{E}\}$
- $\Phi(a) = \underline{A}, \Phi(b) = \underline{B}, \Phi(c) = \underline{C}, \Phi(e) = \underline{E}.$
- $\Psi(\text{on}) = \{(\underline{A}, \underline{B}), (\underline{B}, \underline{C})\}$
- $\Psi(\text{above}) = \{(\underline{A}, \underline{B}), (\underline{B}, \underline{C}), (\underline{A}, \underline{C})\}$
- $\Psi(\text{clear}) = \{\underline{A}, \underline{E}\}$
- $\Psi(\text{ontable}) = \{\underline{C}, \underline{E}\}$

$\forall X \exists Y. (\text{clear}(X) \vee \text{on}(Y, X))$

$X = \underline{A}$

$X = \underline{C}, Y = \underline{B}$

...

$\exists Y \forall X. (\text{clear}(X) \vee \text{on}(Y, X))$

$Y = \underline{A} ?$

$Y = \underline{C} ?$

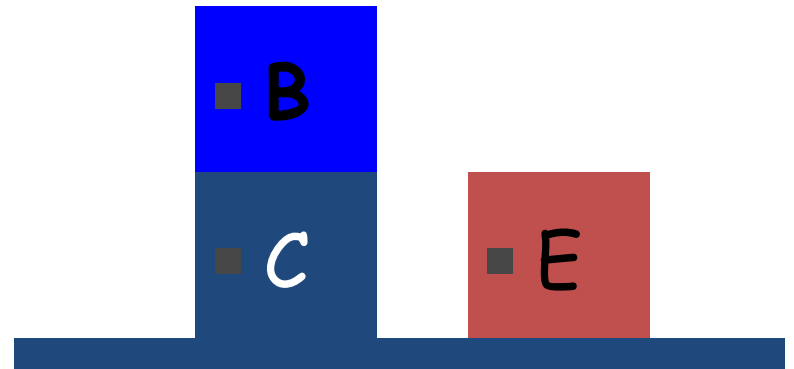
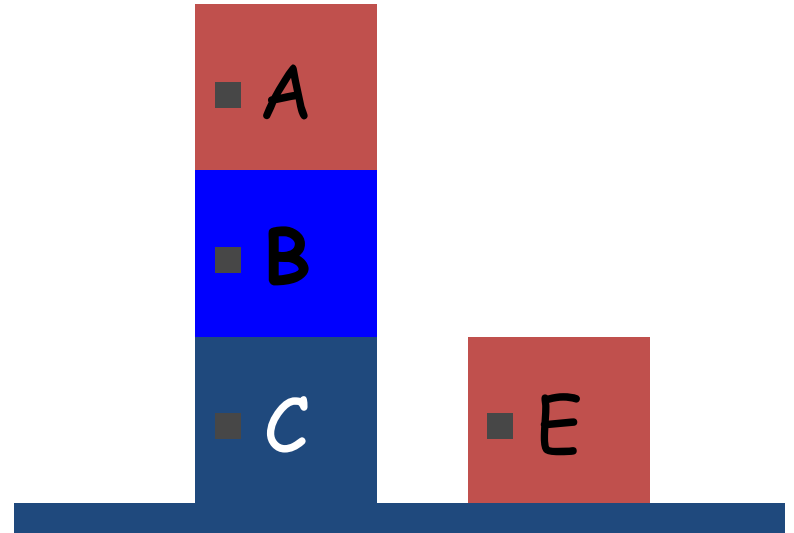
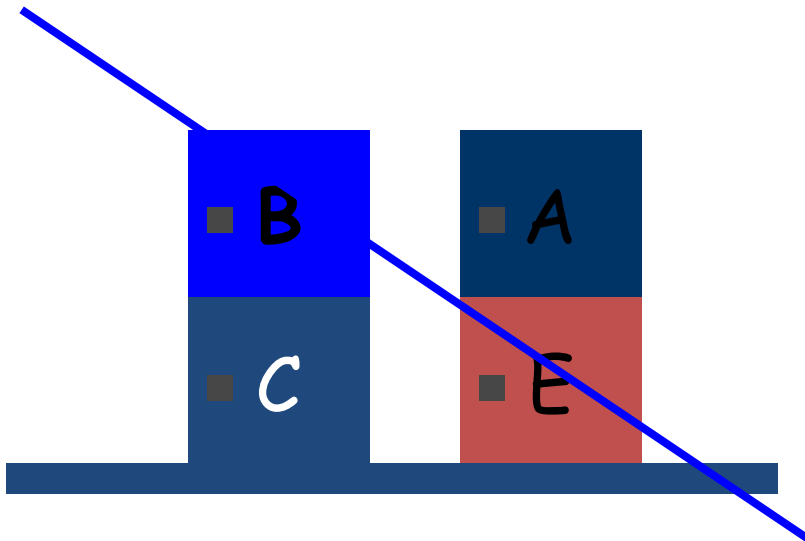
$Y = \underline{E} ?$

$Y = \underline{B} ?$

KB — many models

KB

1. on(b,c)
2. clear(e)



Models

- Let our Knowledge base KB, consist of a set of formulas.
- We say that I is a **model** of KB or that I **satisfies** KB
 - If, every formula $f \in KB$ is true under I
- We write $I \models KB$ if I satisfies KB, and $I \models f$ if f is true under I .

What's Special About Models?

- When we write KB, we intend that the real world (i.e. our set theoretic abstraction of it) is one of its models.
- This means that every statement in KB is **true** in the real world.
- Note however, that not every thing true in the real world need be contained in KB. We might have only incomplete knowledge.

Models support reasoning.

- Suppose formula f is not mentioned in KB, but is true in every model of KB; i.e.,
$$I \models \text{KB} \rightarrow I \models f.$$
- Then we say that f is a **logical consequence** of KB or that KB **entails** f .
- Since the real world is a model of KB, f must be true in the real world.
- This means that entailment is a way of finding new true facts that were not explicitly mentioned in KB.

??? If KB doesn't entail f , is f false in the real world?

Logical Consequence Example

- **elephant(clyde)**
 - the individual denoted by the symbol *clyde* in the set denoted by *elephant* (has the property that it is an *elephant*).
- **teacup(cup)**
 - *cup* is a teacup.
- Note that in both cases a unary predicate specifies a set of individuals. Asserting a unary predicate to be true of a term means that the individual denoted by that term is in the specified set.

Logical Consequence Example

- $\forall X, Y. \text{elephant}(X) \wedge \text{teacup}(Y) \rightarrow \text{largerThan}(X, Y)$
 - For all pairs of individuals if the first is an elephant and the second is a teacup, then the pair of objects are related to each other by the *largerThan* relation.
 - For pairs of individuals who are not elephants and teacups, the formula is immediately true.

Logical Consequence Example

- $\forall X, Y. \text{largerThan}(X, Y) \rightarrow \neg \text{fitsIn}(X, Y)$
 - For all pairs of individuals if X is larger than Y (the pair is in the largerThan relation) then we cannot have that X fits in Y (the pair cannot be in the fitsIn relation).
 - (The relation largerThan has an empty intersection with the fitsIn relation).

Logical Consequences

- $\neg \text{fitsIn}(\text{clyde}, \text{cup})$
- We know $\text{largerThan}(\text{clyde}, \text{teacup})$ from the first implication. Thus we know this from the second implication.

Logical Consequences

fitsIn

\neg fitsIn

largerThan

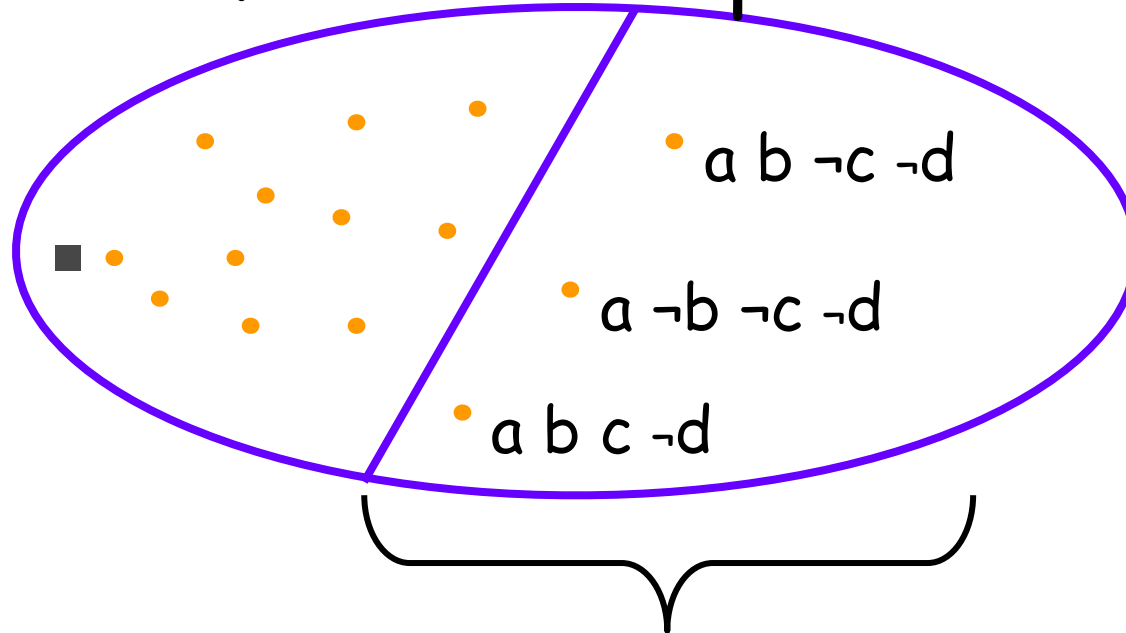
Elephants \times teacups
(clyde , cup)

Logical Consequence Example

- If an interpretation satisfies KB, then the set of pairs *elephant* X *teacup* must be a subset of *largerThan*, which is disjoint from *fitsIn*.
- Therefore, the pair (*clyde*, *cup*) must be in the complement of the set *fitsIn*.
- Hence, $\neg \text{fitsIn}(\text{clyde}, \text{cup})$ must be true in every interpretation that satisfies KB.
- $\neg \text{fitsIn}(\text{clyde}, \text{cup})$ is a logical consequence of KB.

Models Graphically

Set of All Interpretations



Models of KB

a , b , c , and d are atomic formulas

Consequences? $a, c \rightarrow b, b \rightarrow c, d \rightarrow b, \neg b \rightarrow \neg c$

Models and Interpretations

- the more sentences in KB, the fewer models (satisfying interpretations) there are.
- The more you write down (as long as it's all true!), the “closer” you get to the “real world”! Because Each sentence in KB rules out certain unintended interpretations.
- This is called **axiomatizing the domain**

Computing logical consequences

- We want procedures for computing logical consequences that can be implemented in our programs.
- This would allow us to reason with our knowledge
 - Represent the knowledge as logical formulas
 - Apply procedures for generating logical consequences
- These procedures are called **proof procedures**.

Proof Procedures

- Interesting, proof procedures work by simply manipulating formulas. They do not know or care anything about interpretations.
- Nevertheless they respect the semantics of interpretations!
- We will develop a proof procedure for first-order logic called resolution.
 - Resolution is the mechanism used by PROLOG

Properties of Proof Procedures

- Before presenting the details of resolution, we want to look at properties we would like to have in a (any) proof procedure.
- We write $KB \vdash f$ to indicate that f can be proved from KB (the proof procedure used is implicit).

Properties of Proof Procedures

- Soundness

- $KB \vdash f \rightarrow KB \models f$

i.e all conclusions arrived at via the proof procedure are correct: they are logical consequences.

- Completeness

- $KB \models f \rightarrow KB \vdash f$

i.e. every logical consequence can be generated by the proof procedure.

- Note proof procedures for FOL have very high complexity in the worst case. So completeness is not necessarily achievable in practice.